

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: df=pd.read_csv("Salary_Data[1].csv")
```

```
In [4]: df.head(10)
```

Out[4]:

	Age	Gender	Education Level	Job Title	Years of Experience	Salary
0	32.0	Male	Bachelor's	Software Engineer	5.0	90000.0
1	28.0	Female	Master's	Data Analyst	3.0	65000.0
2	45.0	Male	PhD	Senior Manager	15.0	150000.0
3	36.0	Female	Bachelor's	Sales Associate	7.0	60000.0
4	52.0	Male	Master's	Director	20.0	200000.0
5	29.0	Male	Bachelor's	Marketing Analyst	2.0	55000.0
6	42.0	Female	Master's	Product Manager	12.0	120000.0
7	31.0	Male	Bachelor's	Sales Manager	4.0	80000.0
8	26.0	Female	Bachelor's	Marketing Coordinator	1.0	45000.0
9	38.0	Male	PhD	Senior Scientist	10.0	110000.0

```
In [5]: df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0})
df['Education Level'] = df['Education Level'].map({'Bachelor': 0, 'Master': 1, 'PhD': 2})
```

```
In [6]: df = df.dropna()
#dropping the null values
```

```
In [57]: print(df['Education Level'].unique())
#Converting categorical data to numerical data
df['Education Level'] = df['Education Level'].map({'Bachelor': 0, 'Master': 1, 'PhD': 2})
print(df)
```

```
[2.]
      Age  Gender  Education Level  Job Title \
2    45.0    1.0             NaN    Senior Manager
9    38.0    1.0             NaN    Senior Scientist
17   39.0    1.0             NaN    Senior Engineer
28   43.0    0.0             NaN    Senior Consultant
34   46.0    1.0             NaN    Senior Manager
...    ...    ...             ...    ...
6682  28.0    1.0             NaN    Marketing Manager
6685  49.0    0.0             NaN    Senior Product Marketing Manager
6691  36.0    0.0             NaN    Marketing Manager
6696  28.0    1.0             NaN    Sales Representative
6699  49.0    0.0             NaN    Director of Marketing

      Years of Experience  Salary
2                15.0  150000.0
9                10.0  110000.0
17               12.0  115000.0
28               15.0  140000.0
34               20.0  170000.0
...               ...    ...
6682              4.0   55000.0
6685             20.0  200000.0
6691              9.0   95000.0
6696              4.0   55000.0
6699             20.0  200000.0
```

[1368 rows x 6 columns]

```
In [58]: X = df[['Age', 'Gender', 'Years of Experience']].values
y = df['Salary'].values
```

```
In [59]: print(X)
```

```
[[45.  1. 15.]
 [38.  1. 10.]
 [39.  1. 12.]
 ...
 [36.  0.  9.]
 [28.  1.  4.]
 [49.  0. 20.]]
```

```
In [60]: print(y)
```

```
[150000. 110000. 115000. ... 95000. 55000. 200000.]
```

```
In [61]: X = (X - np.mean(X, axis=0)) / np.std(X, axis=0)
```

```
In [62]: X = np.c_[np.ones(X.shape[0]), X]
```

```
In [63]: W = np.zeros(X.shape[1])
```

```
In [64]: def hypothesis(X, W):
         return np.dot(X, W)
```

```
In [65]: def cost_function(X, y, W):
         m = len(y)
         return (1 / (2 * m)) * np.sum((hypothesis(X, W) - y) ** 2)
```

```
In [66]: # Gradient Descent Algorithm
         def gradient_descent(X, y, W, alpha, num_iterations):
             m = len(y)
             cost_history = []

             for i in range(num_iterations):
                 # Updating the weights
                 W = W - (alpha / m) * np.dot(X.T, (hypothesis(X, W) - y))

                 # Calculating cost and store in history
                 cost_history.append(cost_function(X, y, W))

             return W, cost_history
```

```
In [67]: alpha = 0.001
         num_iterations = 1000

         # Training the model
         W, cost_history = gradient_descent(X, y, W, alpha, num_iterations)

         print("Weights:", W)
         print("Cost history:", cost_history[-10:]) # Last 10 cost values
```

```
Weights: [104763.27538946  9784.07848236   969.10878162  10270.9667527 ]
Cost history: [2218459625.0635314, 2214658531.9170775, 2210865076.5842247, 2207079243.6
43146, 2203301017.7034383, 2199530383.406061, 2195767325.423268, 2192011828.4585447, 21
88263877.246542, 2184523456.5530114]
```

```
In [69]: def predict(X, W): return hypothesis(X, W)
```

```
In [70]: predicted_salaries = predict(X, W)
```

```
In [71]: df['Predicted Salary'] = predicted_salaries
```

```
In [72]: print("DataFrame with Predicted Salaries:\n", df.head())
```

```
DataFrame with Predicted Salaries:
   Age  Gender  Education Level  Job Title  Years of Experience  \
2  45.0    1.0         NaN      Senior Manager             15.0
9  38.0    1.0         NaN      Senior Scientist             10.0
17 39.0    1.0         NaN      Senior Engineer             12.0
28 43.0    0.0         NaN      Senior Consultant             15.0
34 46.0    1.0         NaN      Senior Manager             20.0

   Salary  Predicted Salary
2  150000.0    112354.766614
9  110000.0    93925.862301
17 115000.0    99065.181062
28 140000.0   107858.640469
34 170000.0   123342.861048
```

```
In [74]: df.head(10)
```

```
Out[74]:
```

	Age	Gender	Education Level	Job Title	Years of Experience	Salary	Predicted Salary
2	45.0	1.0	NaN	Senior Manager	15.0	150000.0	112354.766614
9	38.0	1.0	NaN	Senior Scientist	10.0	110000.0	93925.862301
17	39.0	1.0	NaN	Senior Engineer	12.0	115000.0	99065.181062
28	43.0	0.0	NaN	Senior Consultant	15.0	140000.0	107858.640469
34	46.0	1.0	NaN	Senior Manager	20.0	170000.0	123342.861048
48	38.0	1.0	NaN	Senior Scientist	11.0	120000.0	95875.454192
57	43.0	1.0	NaN	Senior Engineer	17.0	140000.0	113773.680436
63	47.0	1.0	NaN	Senior Data Scientist	21.0	180000.0	126532.587919
72	45.0	1.0	NaN	Research Director	16.0	190000.0	114304.358505
83	52.0	1.0	NaN	Chief Technology Officer	24.0	250000.0	138582.038491

```
In [ ]:
```